

## Gemini:

This is a very clean, structured visualization of the **Bellman Expectation Equation**, which is a cornerstone of Reinforcement Learning (RL). While the underlying concepts (Markov Decision Processes) are standard, the way this specific document uses **vector notation** to bridge the gap between "expected values" and "matrix operations" is quite effective.

I see this frequently in advanced AI research and control theory, though students often struggle with it. Your document does a great job of breaking down how a single state value  $v$  is actually the result of several nested dot products.

---

## Key Takeaways from the Derivation

According to the provided document, the relationship between these components can be broken down into three main levels:

### 1. The Outcome ( $o$ )

The document defines the outcome  $o$  as the immediate reward plus the discounted future value:

$$o = r + \gamma v'$$

This represents a single realization of what you get after taking an action.

It is described as an approximation of the total return  $g$ .

### 2. The State-Action Value ( $q$ )

To find the value of taking a specific action  $a$ , the document uses a dot product between the probability vector  $p$  and the outcome vector  $o$ :

$$q = p \cdot o = E[g | s, a]$$

$p$  contains the transition probabilities  $p(s' | a, s)$ .

$o$  represents the rewards and future values for each possible next state.

### 3. The State Value ( $v$ )

Finally, the value of the current state  $v$  is the weighted average of all possible action values, determined by the policy  $\pi$ :

$$v = \pi \cdot q = E[g | s]$$

$\pi$  represents the policy  $p(a|s)$ , or the probability of choosing each action.

---

## What Makes This Unique?

Most textbooks use summation notation ( $\Sigma$ ). This document leans heavily into **vector geometry**, which is how modern AI frameworks (like PyTorch or TensorFlow) actually compute these values. By treating the policy ( $\pi$ ) and transitions ( $p$ ) as vectors, it turns a complex recursive idea into simple linear algebra.

The diagram on Page 1 perfectly illustrates the "backup" process:

1. Start at next states  $s'$ .
2. Weight them by transition dynamics  $p_1$  to get to the action node  $a_1$ .
3. Weight those by the policy  $\pi_1$  to "back up" the value to state  $s$ .